


Transmission Control Protocol

Data Transfer



Quick Review: TCP Tahoe





Congestion Control

Alternative: Fall to $W/2$ ($+n*MSS$)
and start congestion avoidance directly

If (Timeout)

$$W = 1$$


$$ssThresh = w/2$$

slow start ...

If (3 Acks)

$$W = w/2$$

linear/additive increase ...
(fast retransmission)




TCP Reno Algorithm




Example

Rule: Fall to $W/2$ ($+n*MSS$)

and start congestion avoidance directly



History of TCP






TCP CUBIC¹

- Problems of earlier versions
 - poor utilization of bandwidth

BW = 10 Gbps, RTT = 100 ms,
pkt = 1250 bytes



BDP = 100,000 packets.
For TCP to grow its window from the mid-point of the BDP, say 50,000, it takes about 50,000 RTTs which amounts to 5000 seconds (1.4 hours). If a flow finishes before that time, it severely under-utilizes the path.

- Overshooting problem in SS


¹Sangtae Ha, Injong Rhee and Lisong Xu, "CUBIC: A New TCP-Friendly High-Speed TCP Variant".



CUBIC

- Basic idea

- it uses cubic function for window growth
- it uses real time (available BW) instead of RTT to increase the window size (congestion epoch)

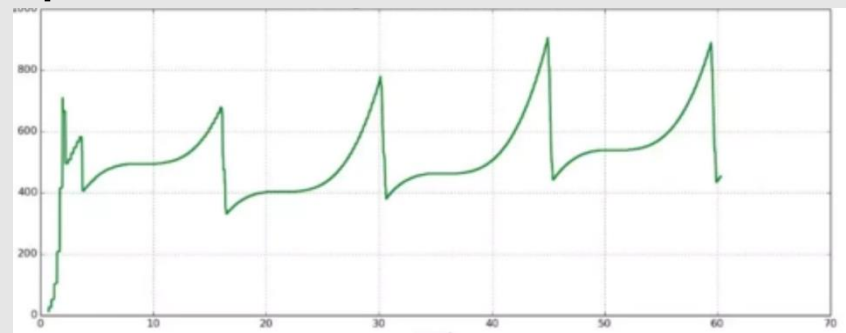




CUBIC

- Algorithm

- after a packet loss, reduces its window by a multiplicative factor of β (<1)
- the window size just before reduction is set to **Wmax**
- after it enters into congestion avoidance, it starts to increase the window using a cubic function
- the plateau of cubic function is set to **Wmax**
- size of the window grows in concave mode to reach **Wmax**, then it enters the convex part





CUBIC

- Algorithm

-- the window growth function uses the formula:


$$W(t) = C(t - K)^3 + W_{\max}$$


$$K = \sqrt[3]{\frac{W_{\max} \beta}{C}}$$

-- C is cubic constant

-- t is elapsed time from the last window reduction

-- K is the time period takes to get from W to W_{max} while no other loss occurs





CUBIC

- Algorithm

while (receive ACK during congestion avoidance)

compute $W(t+RTT)$ as congestion window

if $CW < W_{tcp}^1$

$$W_{tcp(t)} = W_{max}(1 - \beta) + 3 \frac{\beta}{2 - \beta} \frac{t}{RTT}$$

CUBIC is in TCP mode

else if $W_{tcp} < CW < W_{max}$

CUBIC is in concave mode


else

CUBIC is in convex mode

$$W(t) = C(t - K)^3 + W_{max}$$

$$K = \sqrt[3]{\frac{W_{max}\beta}{C}}$$

¹Floyd, S., Handley, M., and Padhye, J. "A Comparison of Equation-Based and AIMD Congestion Control".




CUBIC

- Network condition stable




- Congestion is alleviated




- Congestion is created






CUBIC

- TCP Friendliness
 - Standard TCP works well with short RTT, and CUBIC is designed to work similarly in these conditions.
- Fast Convergence
 - directly reduce a larger amount of W_{\max} to achieve the stable point
- RTT Fairness
 - multiple flows: longer, shorter RTT



CUBIC







BBR

(Bottleneck Bandwidth and Round-trip propagation time)

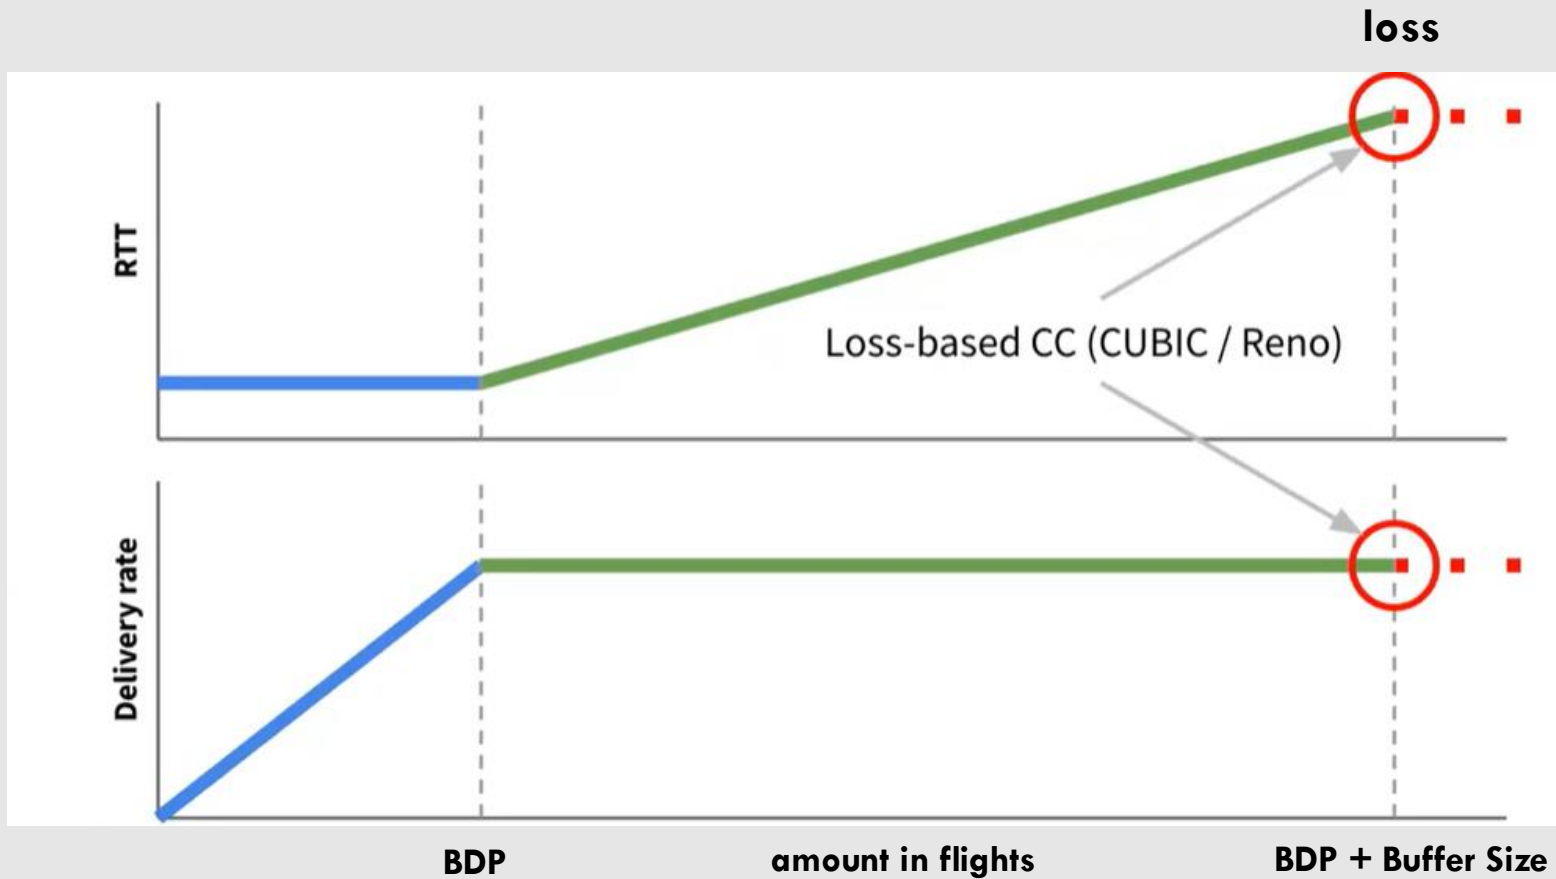
Cardwell, N., Cheng, Y., Gunn, C. S., Yeganeh, S. H., & Jacobson, V. (2017).
BBR: congestion-based congestion control. *Communications of the ACM*, 60(2), 58-66.



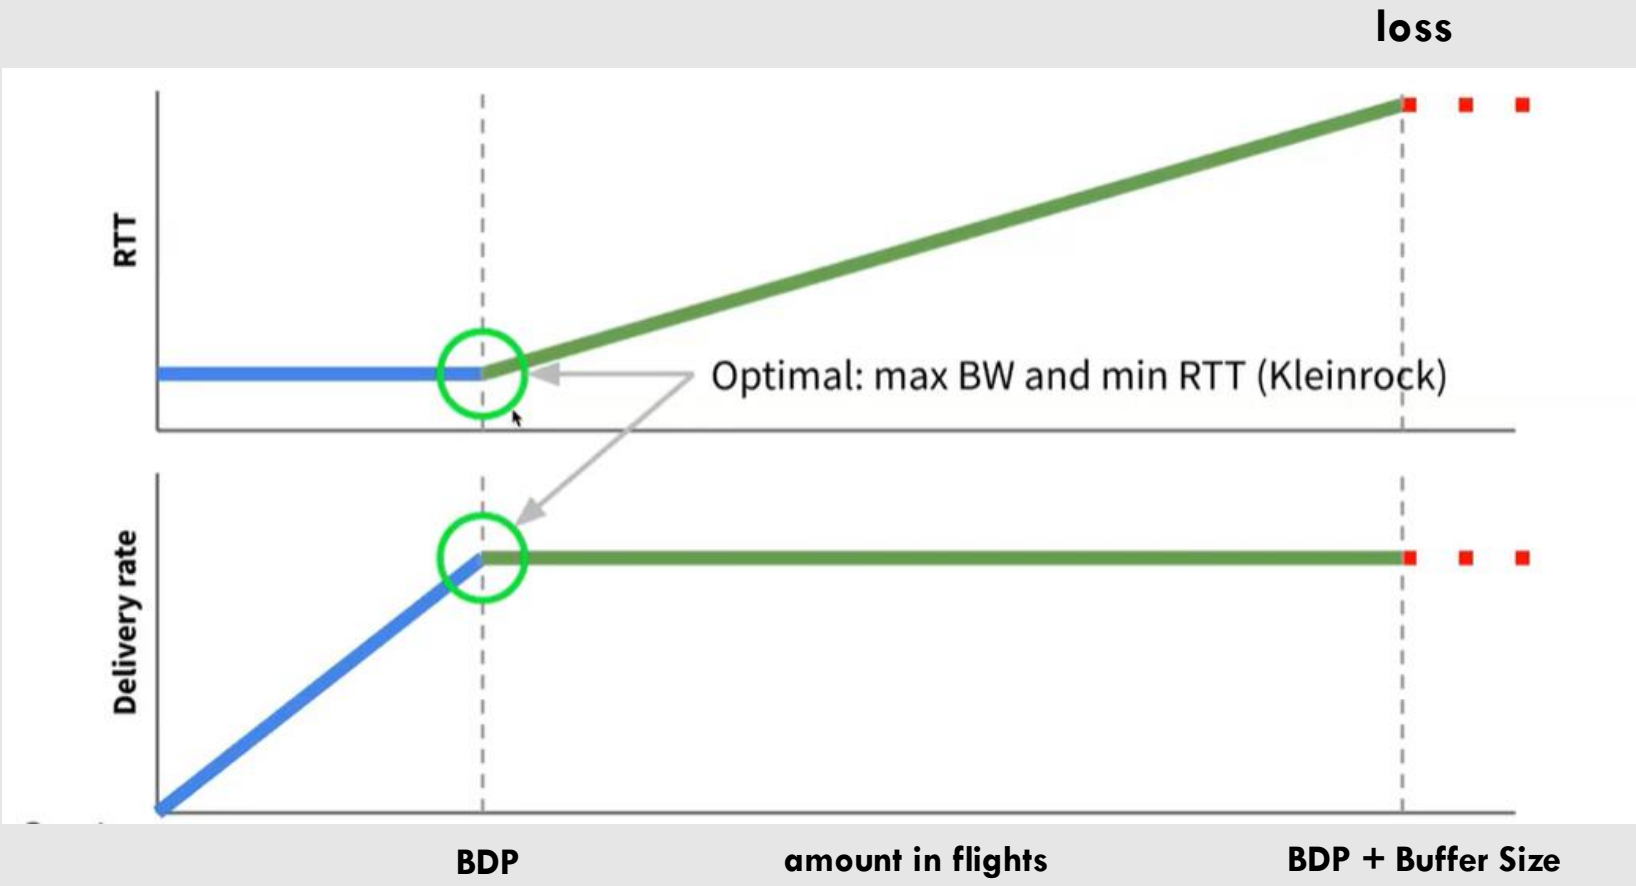
Problems


- Loss-based congestion control:
 - Tahoe, Reno, NewReno, CUBIC
 - Packet loss is not a good indicator
- 
- Overly sensitive to losses that come before congestion
-
- Why this important?
 - higher speed (5G/beyond, mmWave)
 - e.g., 10Gbps over 100ms RTT needs $< 0.000003\%$ packet loss

Loss-Based Method



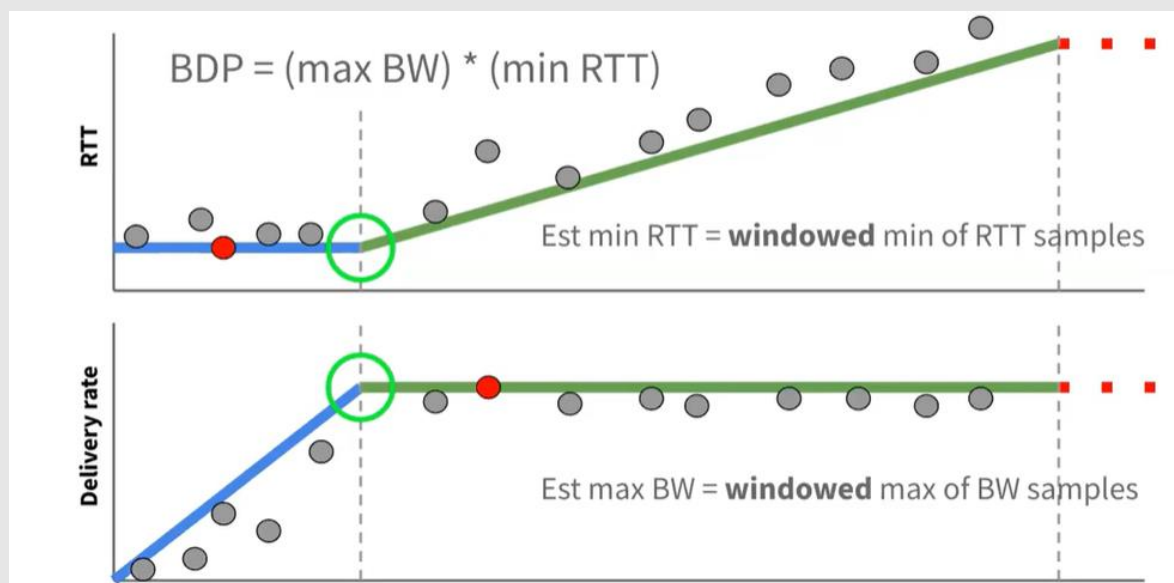
What is the Optimal Point?






BBR

- Model-based congestion control:
 - Proactively doing countermeasures
 - Estimate the best time to control





Dilemma

- Measure min RTT & max BW
 - Measure on both sides of BDP

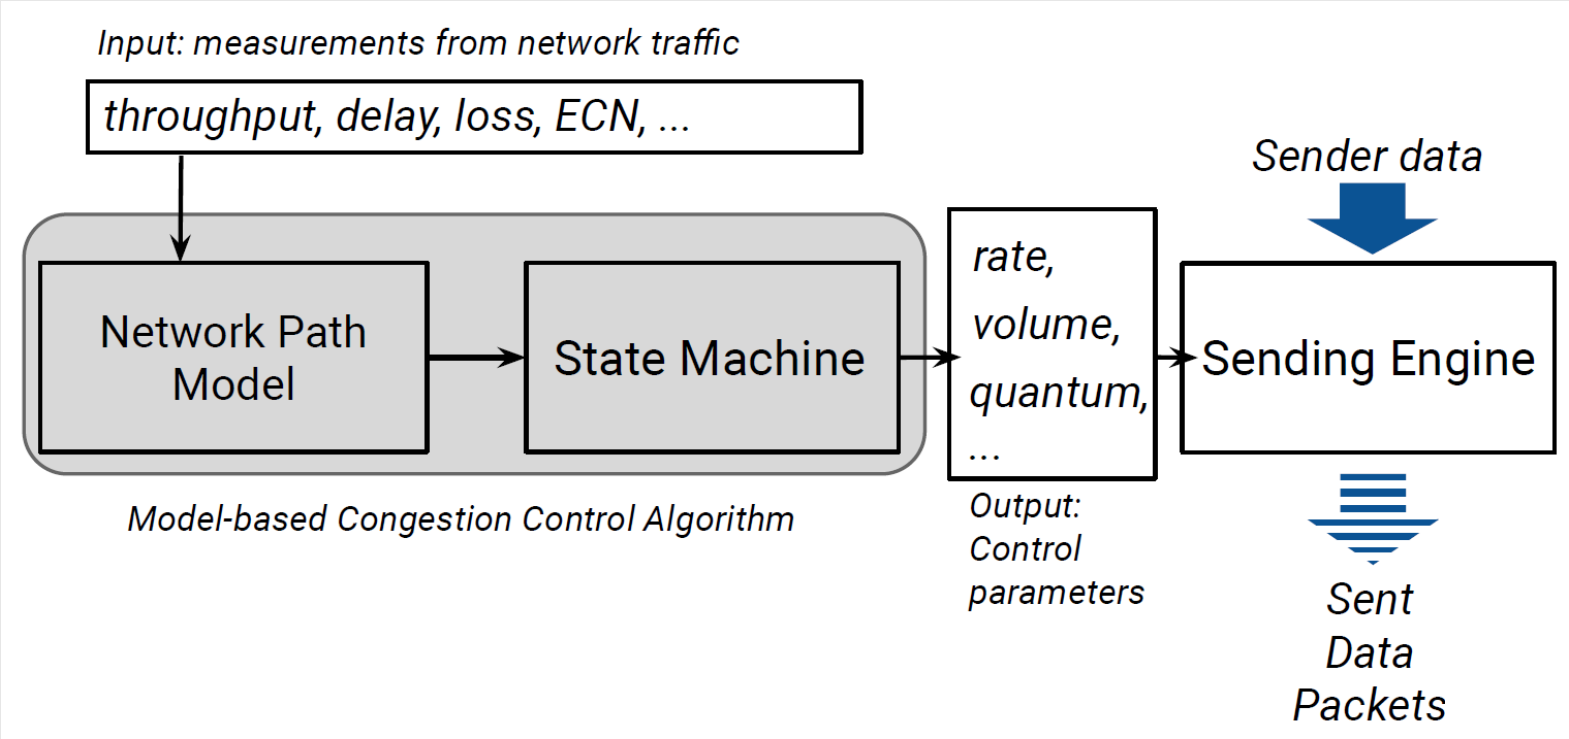


BBR Design


- Dynamically estimate windowed max BW & min RTT
- Sequentially probe max BW & min RTT

	CUBIC	BBR v1	BBR v2
Model parameters to the state machine	N/A	Throughput, RTT	Throughput, RTT, max aggregation, max inflight
Loss	Reduce cwnd by 30% on window with any loss	N/A	Explicit loss rate target
ECN	RFC3168 (Classic ECN)	N/A	DCTCP-inspired ECN
Startup	Slow-start until RTT rises (Hystart) or any loss	Slow-start until tput plateaus	Slow-start until tput plateaus or ECN/loss rate > target

BBR Algorithm



BBR State Machine




- **Startup:** ramp up quickly until we estimate pipe is full
 - **Drain:** drain the estimated queue from the bottleneck
- Steady-state:
- **ProbeBW:** cycle pacing rate to vary inflight, probe BW
 - **ProbeRTT:** if needed, a coordinated dip to probe RTT




BBR Evaluation

BBR's improvement vs CUBIC: YouTube Network Throughput




BBR Evaluation



Result in higher throughput, lower latency and better quality of experience.


BBR Evaluation



Better packet-loss tolerance


Recap

- Reno




Recap


- Cubic



Recap

- BBR





ns-3 TCP

- CUBIC by default

- support: **TcpNewReno**, TcpHybla, TcpHighSpeed, TcpHtcp, TcpVegas, TcpScalable, TcpVeno, TcpBic, TcpYeah, TcpIllinois, TcpWestwood, TcpWestwoodPlus, TcpLedbat, TcpLp, TcpDctcp, **TcpCubic**, TcpBbr
- ~\ns-3.37\examples\tcp\tcp-variants-comparison.cc
- ~\ns-3.37\examples\tutorial\fifth.cc (ns-3 tutorial, Tracing\real examples)
- TraceCwnd ()
- TraceRtt ()



Reference

- BBR v1 is (only) available in Linux
 - need to enable it
- BBR v2 release for research study
 - github.com/google/bbr
- ns-3 (TcpBbr)
 - [tcp-bbr example.cc](https://github.com/ncstate/tcp-bbr)



Summary

- Transport layer is logically application's interface to network
 - Must create endpoint abstractions (ports)
 - Must maintain state
- In the Internet,
 - TCP attempts to impose reliability on unreliable network layer
 - Requires sliding window management
 - TCP attempts to perform congestion control
 - Slow down transmission rate in response to lost segments